

# Case study of sparse matrix intrinsic in LLVM

Eric Feng & Chris Yoon (2024)



# Sparse Matrices

- Compact representation of matrices with lots of zero elements
- Generally stores only the nonzero elements
- Sparse matrix operations only operate on nonzero elements
- Given the right circumstances, huge memory and performance gains!

$$A = \begin{pmatrix} a_{0,0} & 0 & 0 & a_{0,3} \\ 0 & 0 & 0 & 0 \\ a_{2,0} & 0 & 0 & 0 \end{pmatrix}$$

pointers[1]:

indices[1]:

values:

0	2	2	3
0	3	0	
$a_{0,0}$	$a_{0,3}$	$a_{2,0}$	

# LLVM Intrinsics

- Operations that can be matched in the compiler backend to produce the corresponding machine code

```
71     %vec.gep25 = getelementptr float, ptr %r, i64 2
72     %col.load26 = load <2 x float>, ptr %vec.gep25, align 4
73     %18 = load ptr, ptr %pa.addr, align 8
74     store <2 x float> %col.load24, ptr %18, align 4
75     %vec.gep27 = getelementptr float, ptr %18, i64 4
76     store <2 x float> %col.load26, ptr %vec.gep27, align 4
77     ret void
78 }
79
80 ; Function Attrs: nocallback nofree nosync nounwind willreturn memory(argmem: read)
81 declare <4 x float> @llvm.matrix.column.major.load.v4f32.i64(ptr nocapture, i64, i1 immarg, i32 immarg, i32 immarg) #1
82
83 ; Function Attrs: nocallback nofree nosync nounwind speculatable willreturn memory(none)
84 declare <4 x float> @llvm.matrix.multiply.v4f32.v4f32.v4f32(<4 x float>, <4 x float>, i32 immarg, i32 immarg, i32 immarg) #2
85
86 ; Function Attrs: nocallback nofree nosync nounwind willreturn memory(argmem: write)
87 declare void @llvm.matrix.column.major.store.v4f32.i64(<4 x float>, ptr nocapture writeonly, i64, i1 immarg, i32 immarg, i32 immarg) #3
88
```

In LLVM IR emitted from C programs

# LLVM Intrinsics

- Operations that can be matched in the compiler backend to produce the corresponding machine code

```
2009    /// Lowers llvm.matrix.multiply.
2010    void LowerMultiply(CallInst *MatMul) {
2011        IRBuilder<> Builder(MatMul);
2012        auto *EltType = cast<VectorType>(MatMul-&gtgetType())->getElementType();
2013        ShapeInfo LShape(MatMul->getArgOperand(2), MatMul->getArgOperand(3));
2014        ShapeInfo RShape(MatMul->getArgOperand(3), MatMul->getArgOperand(4));
2015
2016        const MatrixTy &Lhs = getMatrix(MatMul->getArgOperand(0), LShape, Builder);
2017        const MatrixTy &Rhs = getMatrix(MatMul->getArgOperand(1), RShape, Builder);
2018        assert(Lhs.getElementType() == Rhs.getElementType() &&
2019            "Matrix multiply argument element types do not match.");
2020
2021        const unsigned R = LShape.NumRows;
2022        const unsigned C = RShape.NumColumns;
2023        assert(LShape.NumColumns == RShape.NumRows);
2024
2025        // Initialize the output
2026        MatrixTy Result(R, C, EltType);
2027        assert(Lhs.getElementType() == Result.getElementType() &&
2028            "Matrix multiply result element type does not match arguments.");
2029
2030        emitMatrixMultiply(Result, Lhs, Rhs, Builder, false, false,
2031            getFastMathFlags(MatMul));
2032        finalizeLowering(MatMul, Result, Builder);
2033    }
2034
```

Generating machine code for the matrix multiplication intrinsics

# Sparse matrix intrinsics in LLVM

## Lowering pipeline:

1. Specify tablegen for desired intrinsic (e.g load, store, multiply)
2. Extend pass manager with option to lower sparse matrix intrinsics
3. Traverse IR and replace all calls to desired intrinsics with custom lowering scheme

## Internal representation:

- Store CSC representation as a contiguous flat vector
  - `[col_ptrs + row_indices + values]`
- Recover CSC for sparse matrix operations

# Limitations of LLVM

- Intrinsics are limited to zero to one return values (LLVM first class types); so we must store CSC representation as a flat vector
- But recovering the flat vector requires us to know values (such as number of nonzeros) that are only available at runtime
- Unable to read the values of `col_ptrs` and `row_indices` at compile time; which are necessary for operations such as SpGEMM, SpMM etc

**Thank you**